T29 – New Employee Onboarding Buddy

Project Report

Working Bot / AI Studio App: https://ai.studio/apps/c81b50a8-78d4-4994-b2c3-ffe92c3cc52e?fullscreenApplet=true

1. Project Objective

The objective of this project was to build an AI-powered new-employee onboarding buddy that provides quick, accurate and source-grounded guidance during the onboarding journey. The bot was designed around the provided onboarding FAQ and the 30-60-90 day checklist. It supports routine questions about Day 1 procedures, IT, benefits, policies, leave, work practices and travel, while also helping employees understand their onboarding milestones.

2. Knowledge Base and Bot Design

The solution uses the supplied 20-question onboarding knowledge base and seven-item 30-60-90 checklist. The checklist covers Day 1, Week 1, Day 30, Day 60 and Day 90 activities, including the responsible owner for each activity. The bot was instructed to use only the supplied information, avoid unsupported assumptions, and explicitly state when information is not available.

Key features include FAQ answering, stage-specific onboarding guidance, a “What should I do this week?” feature for Week 1, and complete stage-specific checklist responses. The interface was built as a simple web-based chatbot with suggested questions and dedicated views for the 30-60-90 tasks and FAQ guidance.

3. Functional Testing

Test	Result
Day 1 guidance	Passed
Week 1 – both tasks	Passed
Laptop FAQ	Passed
Health insurance FAQ	Passed
Unsupported bonus question	Passed
Ambiguous question handling	Passed after improvement
Day 30 checklist	Passed
Day 60 checklist	Passed
Day 90 checklist	Passed
Confidential AI-use question	Passed

A notable defect was found during testing: the initial response to the broad question “What do I need to do?” incorrectly assumed a Day 1 document query. This was corrected by adding an explicit ambiguity rule requiring the user to specify their onboarding stage. The revised behavior was then retested successfully.

4. Prompting and Iteration

Three documented improvement rounds were completed during development. The table below records the problem, the change made, and the observed outcome.

Round	Issue	Improvement	Outcome
1	Risk of unsupported or hallucinated information	Added strict source-grounding and an explicit rule to state when information is unavailable.	Unsupported bonus question was handled correctly without inventing an answer.
2	Ambiguous broad onboarding question	Added a clarification rule asking for Day 1, Week 1, Day 30, Day 60 or Day 90.	Ambiguous query now triggers clarification instead of assuming a stage.
3	Incomplete stage-specific checklist response	Required the complete checklist for the selected stage, including the task owner.	Day 60 response returned the complete task and owner.

The project maintained an AI Use Log with 18 documented entries including the supplied example. The log records prompts, outcomes, changes and whether an output was kept or rejected, providing evidence of iterative development.

5. User Testing

The finished bot was tested by six classmates using realistic onboarding scenarios covering Day 1, Week 1 and Benefits/FAQ questions. Testers rated accuracy, clarity, usefulness, ease of use, response speed and instruction clarity, and provided qualitative feedback.

Metric	Average / Finding
Accuracy	4.50 / 5
Clarity	4.50 / 5
Usefulness	4.67 / 5
Ease of use	4.50 / 5
Response speed	5.00 / 5 (5 responses)
Instruction clarity	4.67 / 5
Incorrect/confusing answers reported	0 of 6
Technical issues reported	0 of 6

Overall, the feedback was strongly positive. One tester gave lower 3/5 ratings for accuracy, clarity and ease of use but did not identify a factual error or confusing response. This was treated as a minor user-experience improvement opportunity rather than evidence of an incorrect answer.

6. Business Value and Recommendation

The onboarding buddy can reduce repetitive questions directed to HR, IT and managers by giving new employees immediate access to common onboarding information. It also gives employees a structured view of what is expected at each stage of the first 90 days. The owner field on checklist tasks makes responsibilities clear and can help new joiners understand whom to approach for specific activities.

Recommendation: Use the onboarding buddy as a first-line self-service onboarding assistant for routine questions, while retaining HR, IT and managers as the escalation points for issues that are not covered by the approved knowledge base. Before organizational deployment, the knowledge base should be maintained by an authorized owner whenever policies, benefits or procedures change.

7. Responsible AI

The bot was deliberately constrained to the supplied knowledge. It does not guess unsupported information, and it was tested using an unknown-information question about the annual performance bonus. It correctly stated that the information was not available. The bot also follows the provided AI-use policy: only company-approved AI tools may be used with work data and confidential information should not be pasted into public AI tools.

8. Limitations and Reflection

The current user test involved six classmates and therefore provides useful usability evidence but is not a full organizational pilot. The bot's knowledge is limited to the supplied FAQ and checklist, so it cannot answer policy changes or company-specific questions that are outside that source. The testing process also showed the importance of explicit ambiguity handling: even a simple question can be interpreted incorrectly if the user's onboarding stage is not known.

The main learning from the project was that building the chatbot is only one part of an effective AI solution. Source grounding, systematic testing, iteration, user feedback and responsible-AI controls were necessary to make the bot reliable and suitable for the intended onboarding use case.

9. Conclusion

The New Employee Onboarding Buddy successfully provides source-grounded onboarding assistance and 30-60-90 day task guidance. Functional testing confirmed the major features, user testing produced strong ratings, and multiple prompting iterations improved reliability. The completed workbook, AI Use Log, user-testing evidence and this report together document the development process and provide a basis for further improvement and controlled deployment.